

Chapter 3: Introduction to Processes

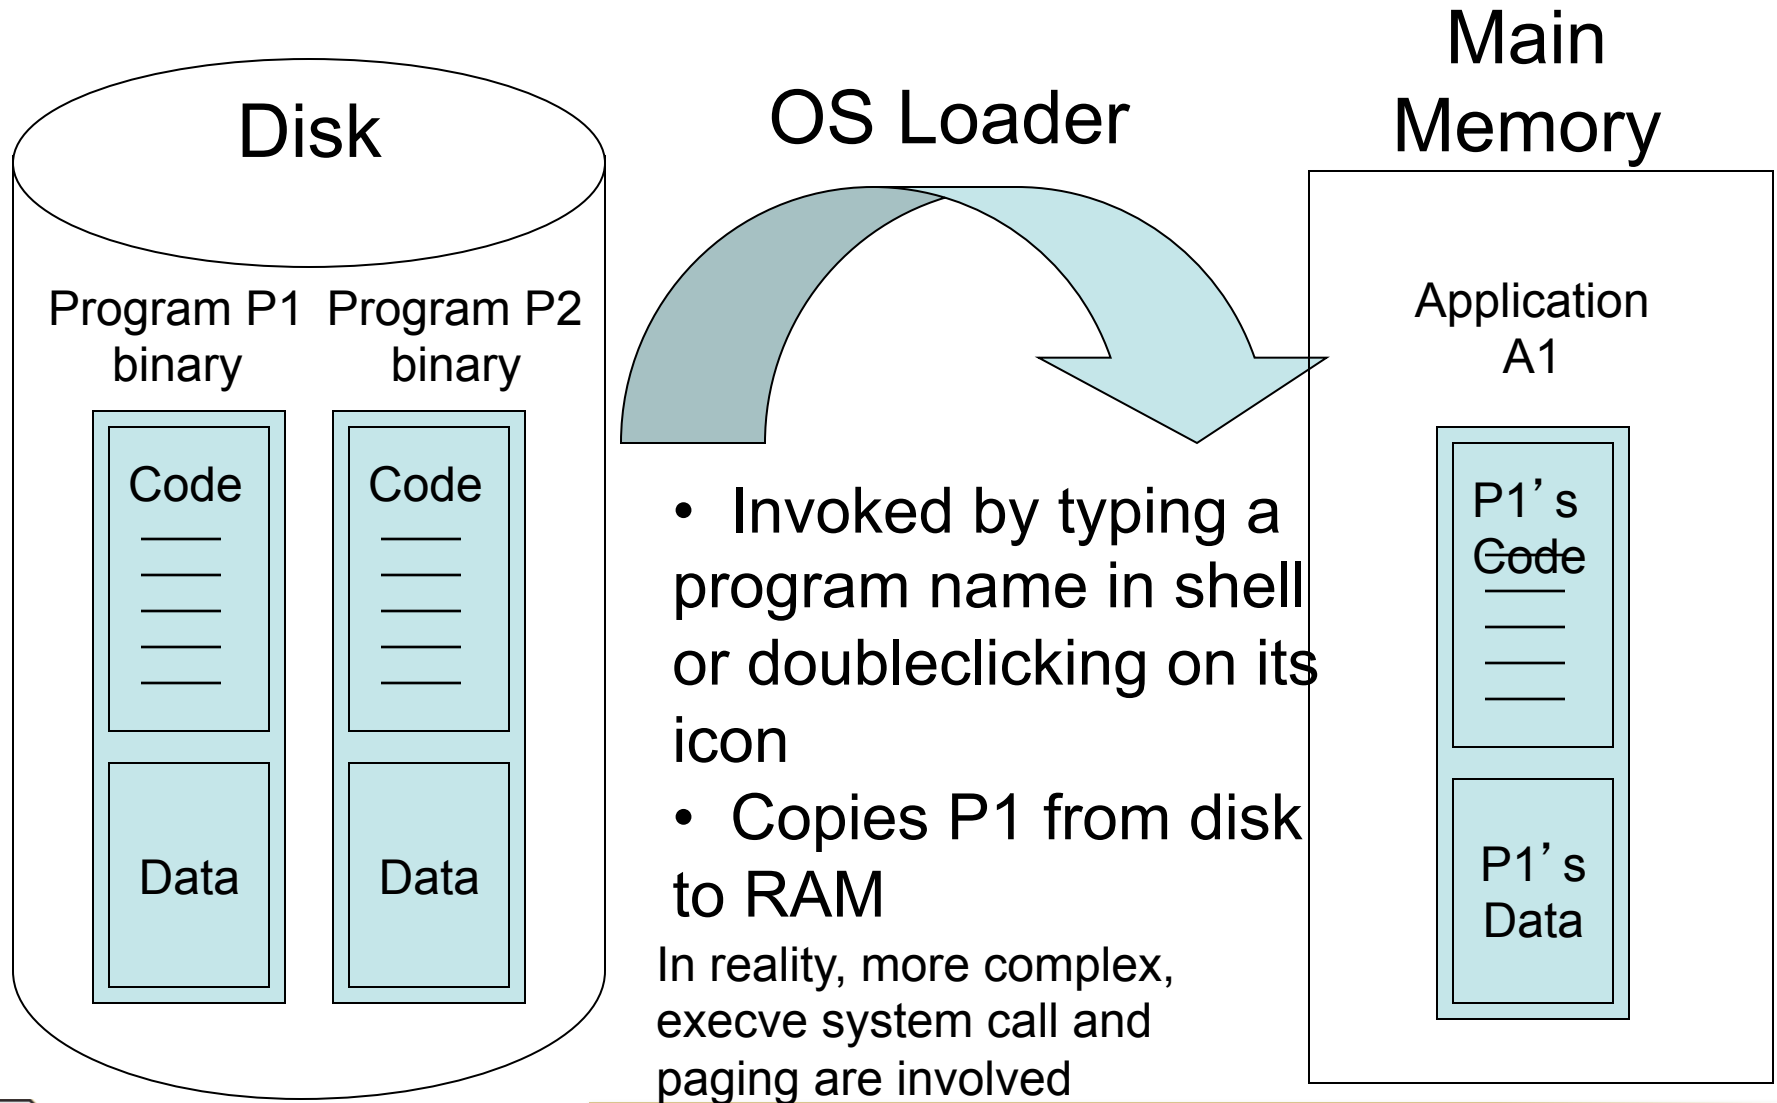
CSCI 3753 Operating Systems

Prof. Rick Han

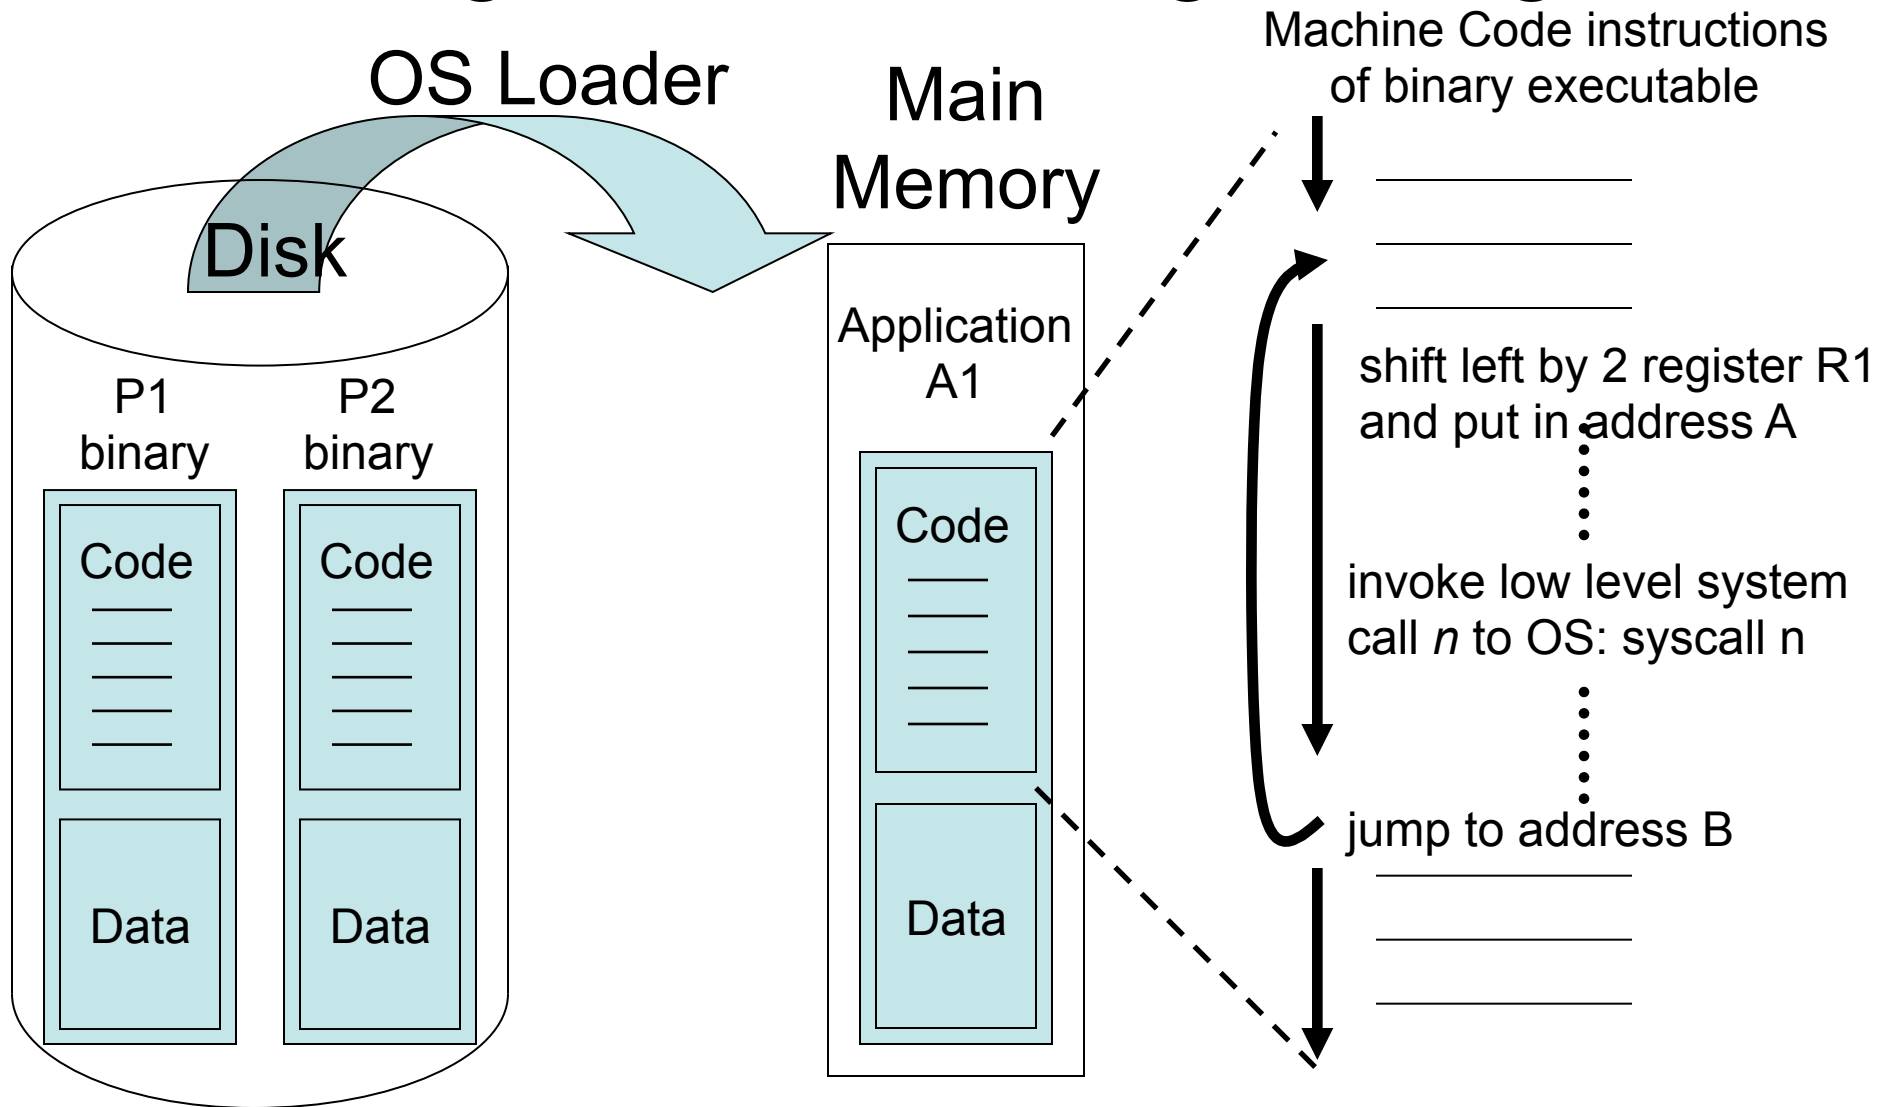
University of Colorado at Boulder



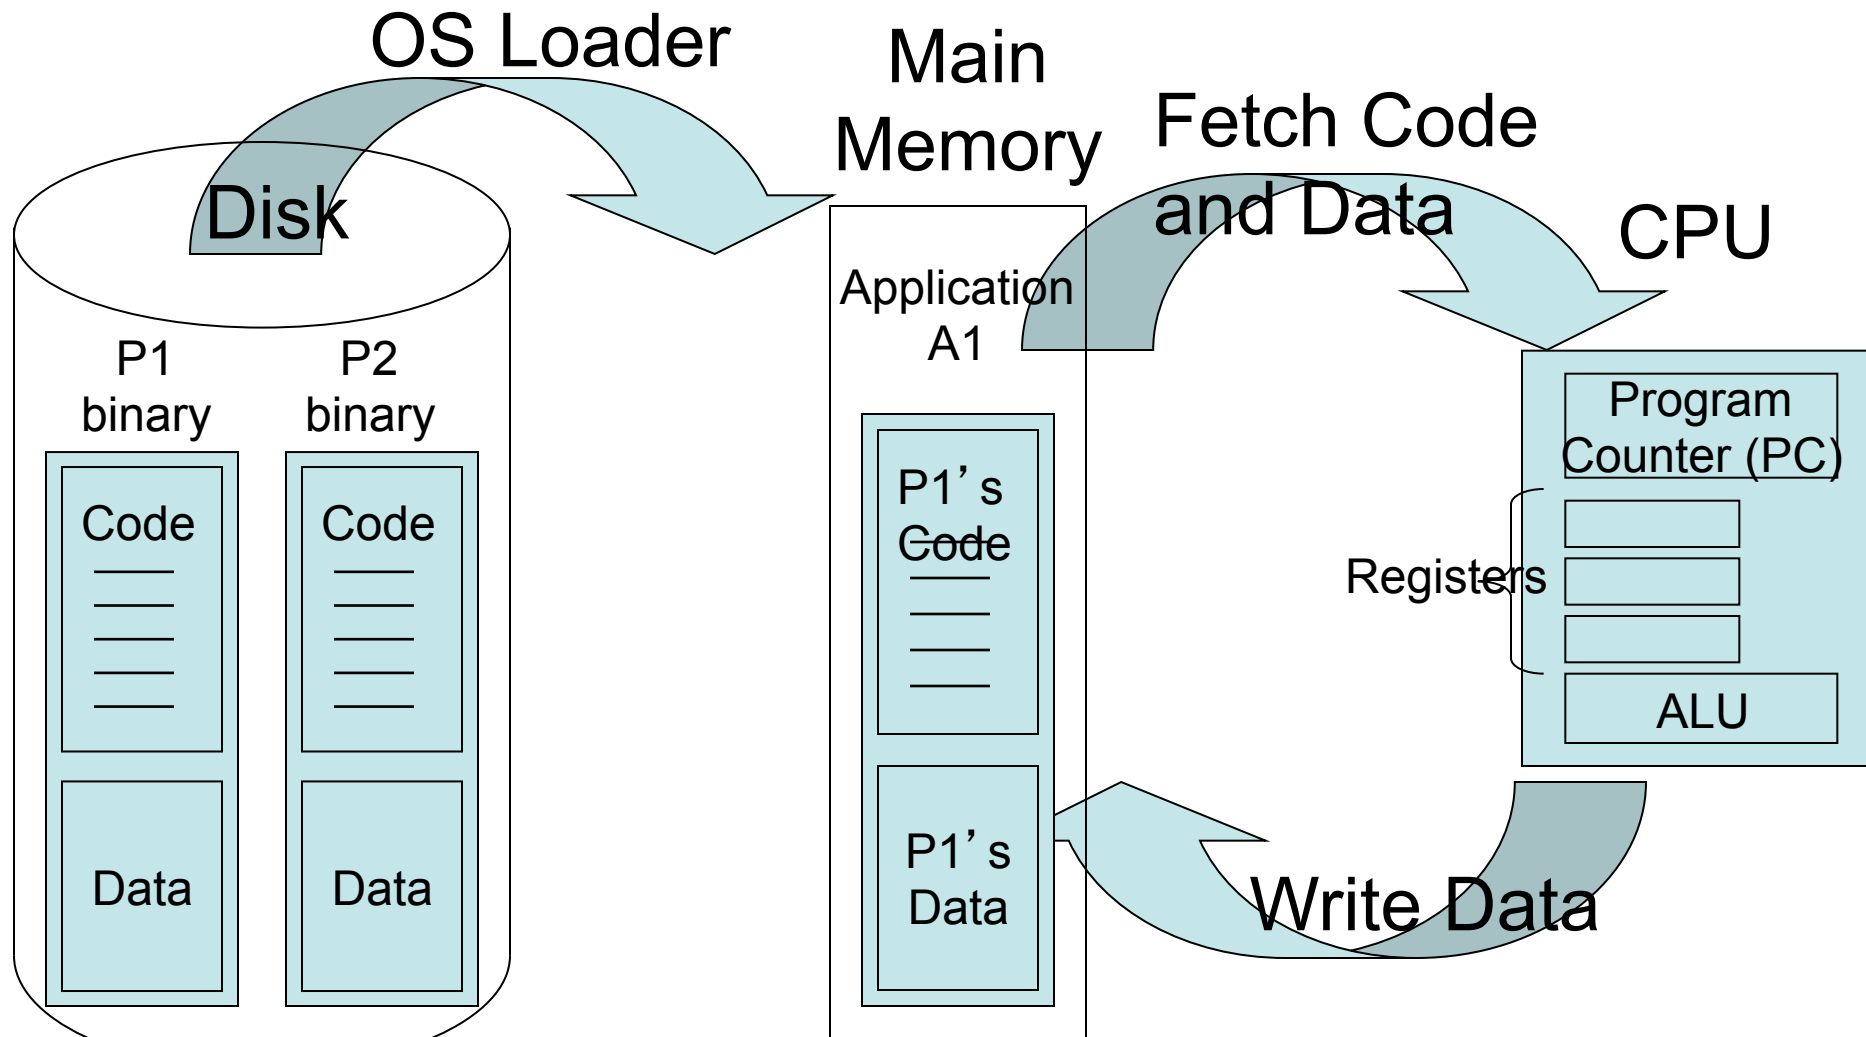
Loading a Program into Memory



Loading and Executing a Program

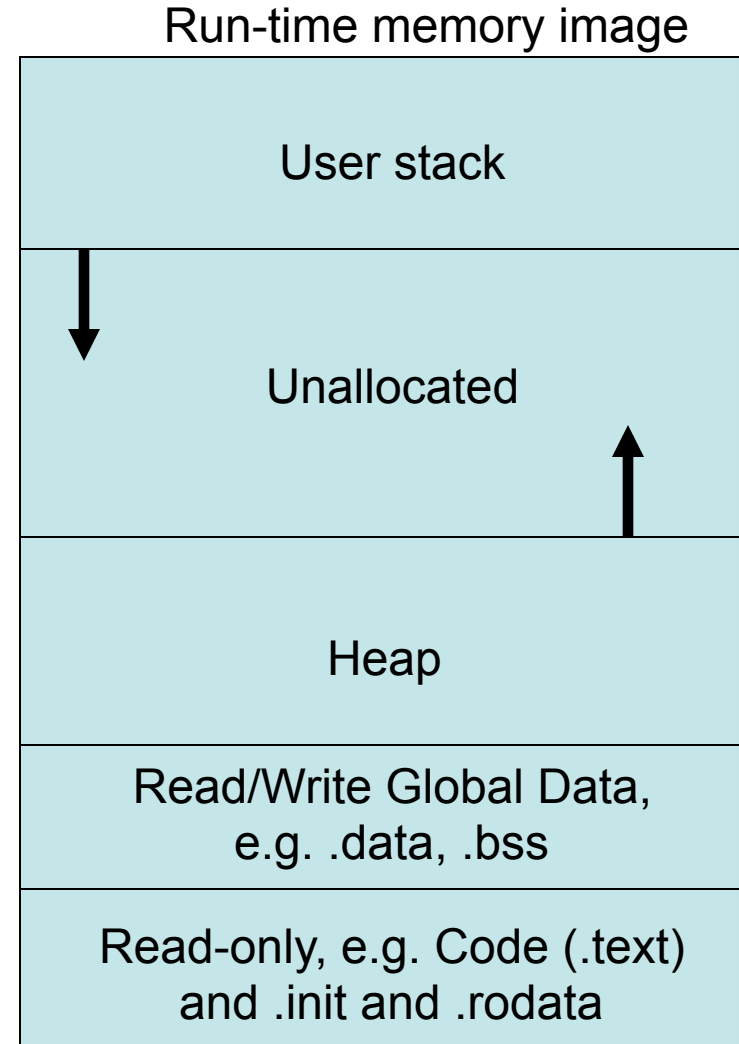


Loading and Executing a Program



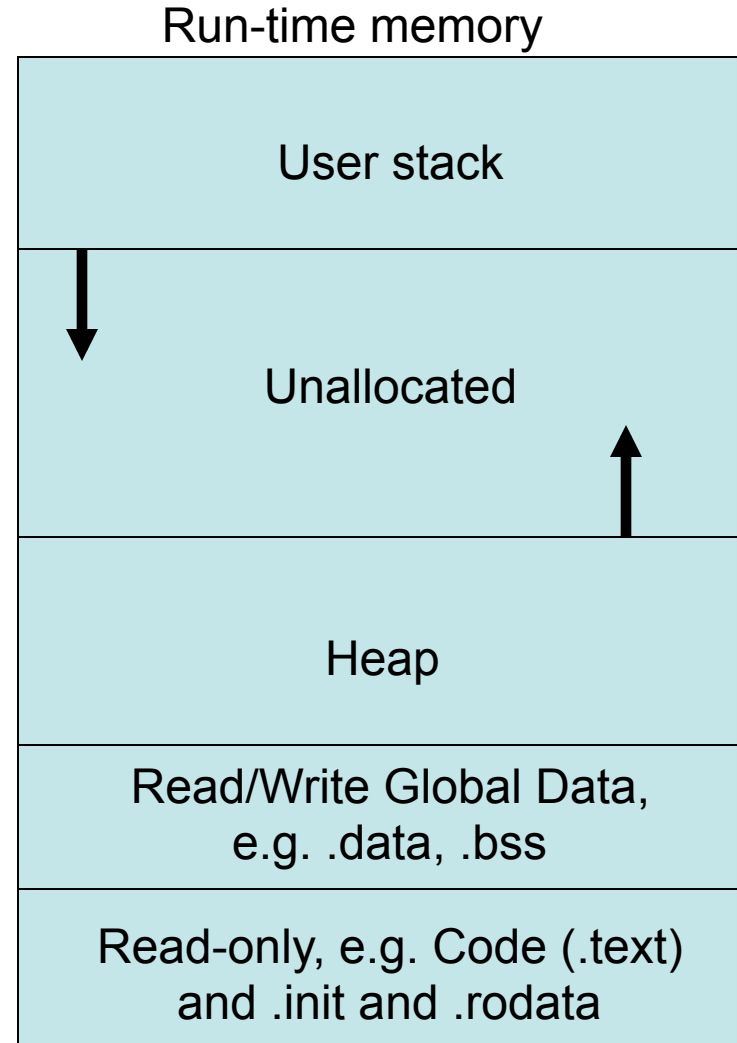
Loading Executable Object Files

- When a program is loaded into RAM, it becomes an actively executing application
- The OS allocates a stack and heap to the app in addition to code and global data.
 - A call stack is for local variables
 - A heap is for dynamic variables, e.g. malloc()
 - Usually, stack grows downward from high memory, heap grows upward from low memory, but this is architecture-specific



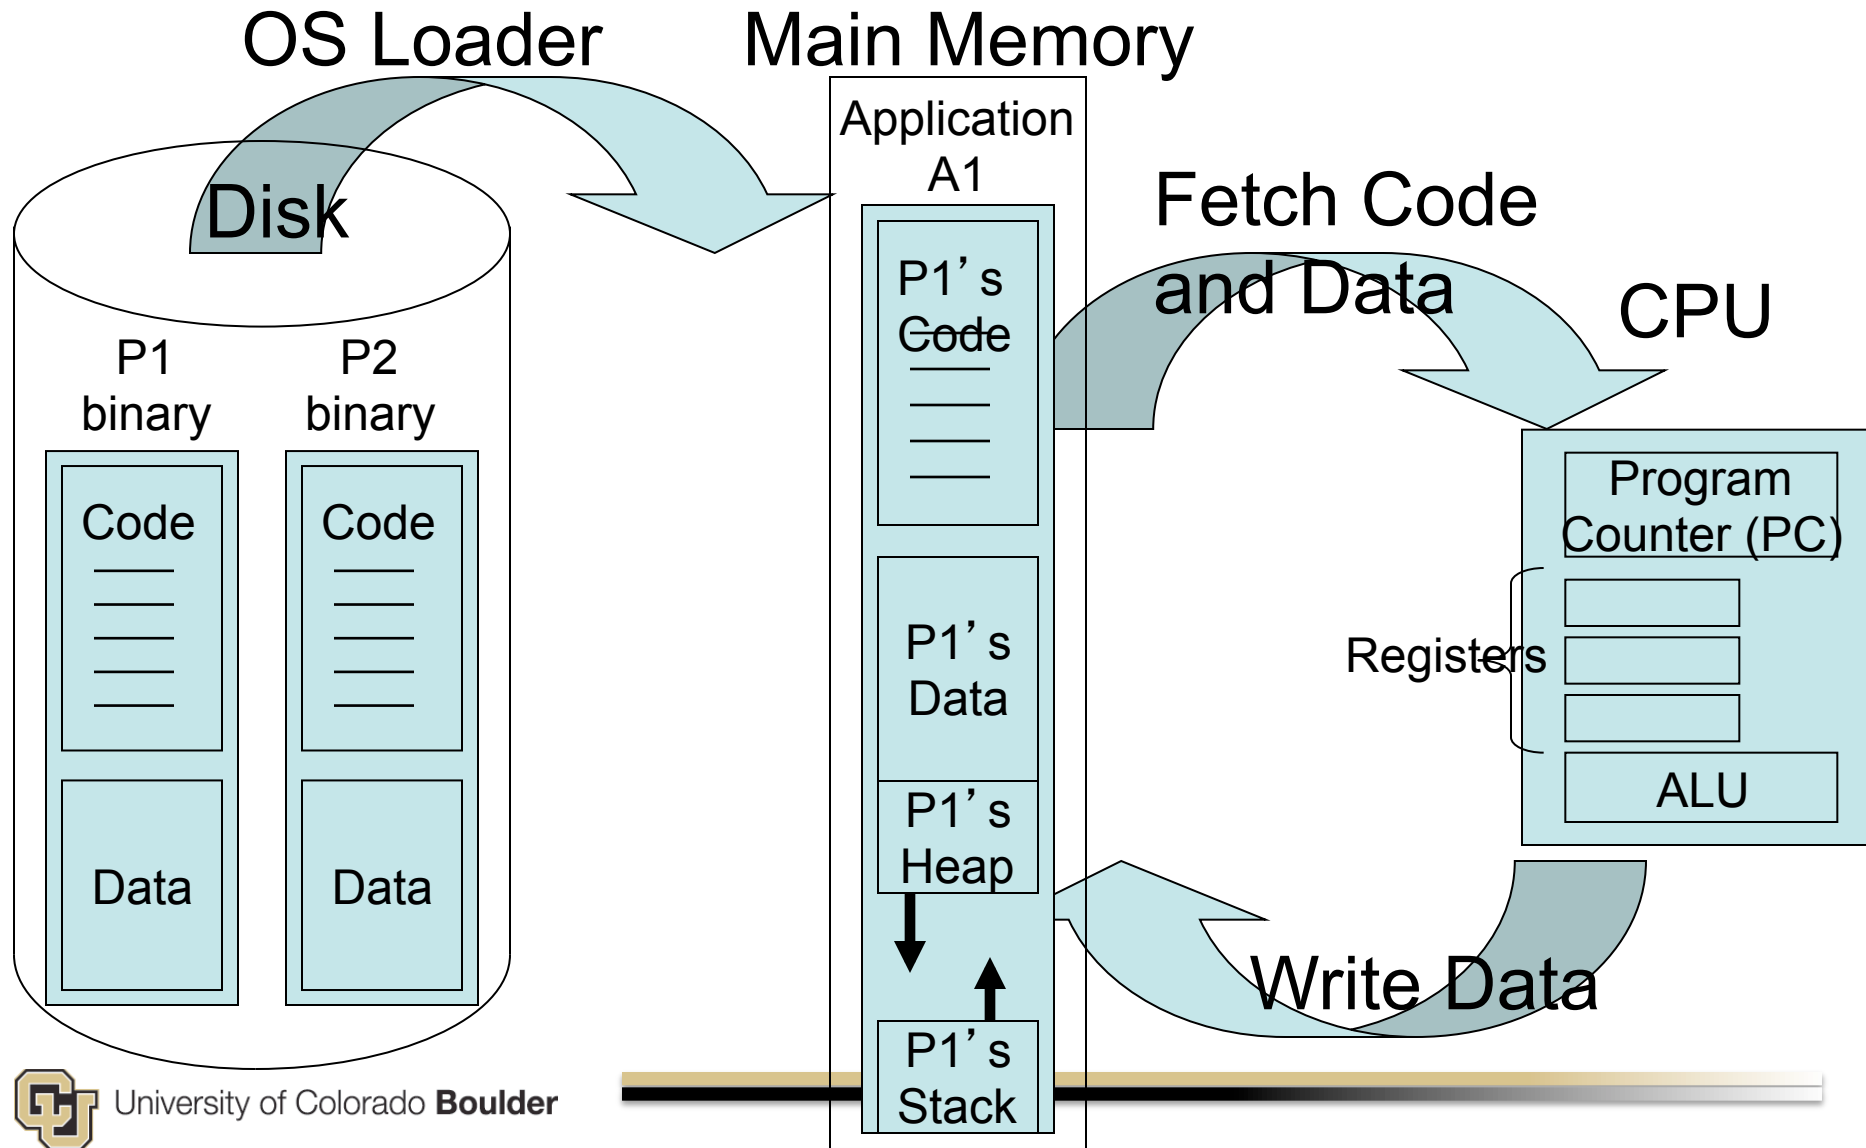
Running Executable Object Files

- Stack contains local variables
 - As `main()` calls function `f1`, we allocate `f1`'s local variables on the stack
 - If `f1` calls `f2`, we allocate `f2`'s variables on the stack below `f1`'s, thereby growing the stack, etc...
 - When `f2` is done, we deallocate `f2`'s local variables, popping them off the stack, and return to `f1`
- Stack dynamically expands and contracts as program runs and different levels of nested functions are called
- Heap contains run-time variables/buffers
 - Obtained from `malloc()`
 - Program should `free()` the `malloc`'ed memory
- Heap can also expand and contract during program execution



Loading and Executing a Program

– a more complete picture

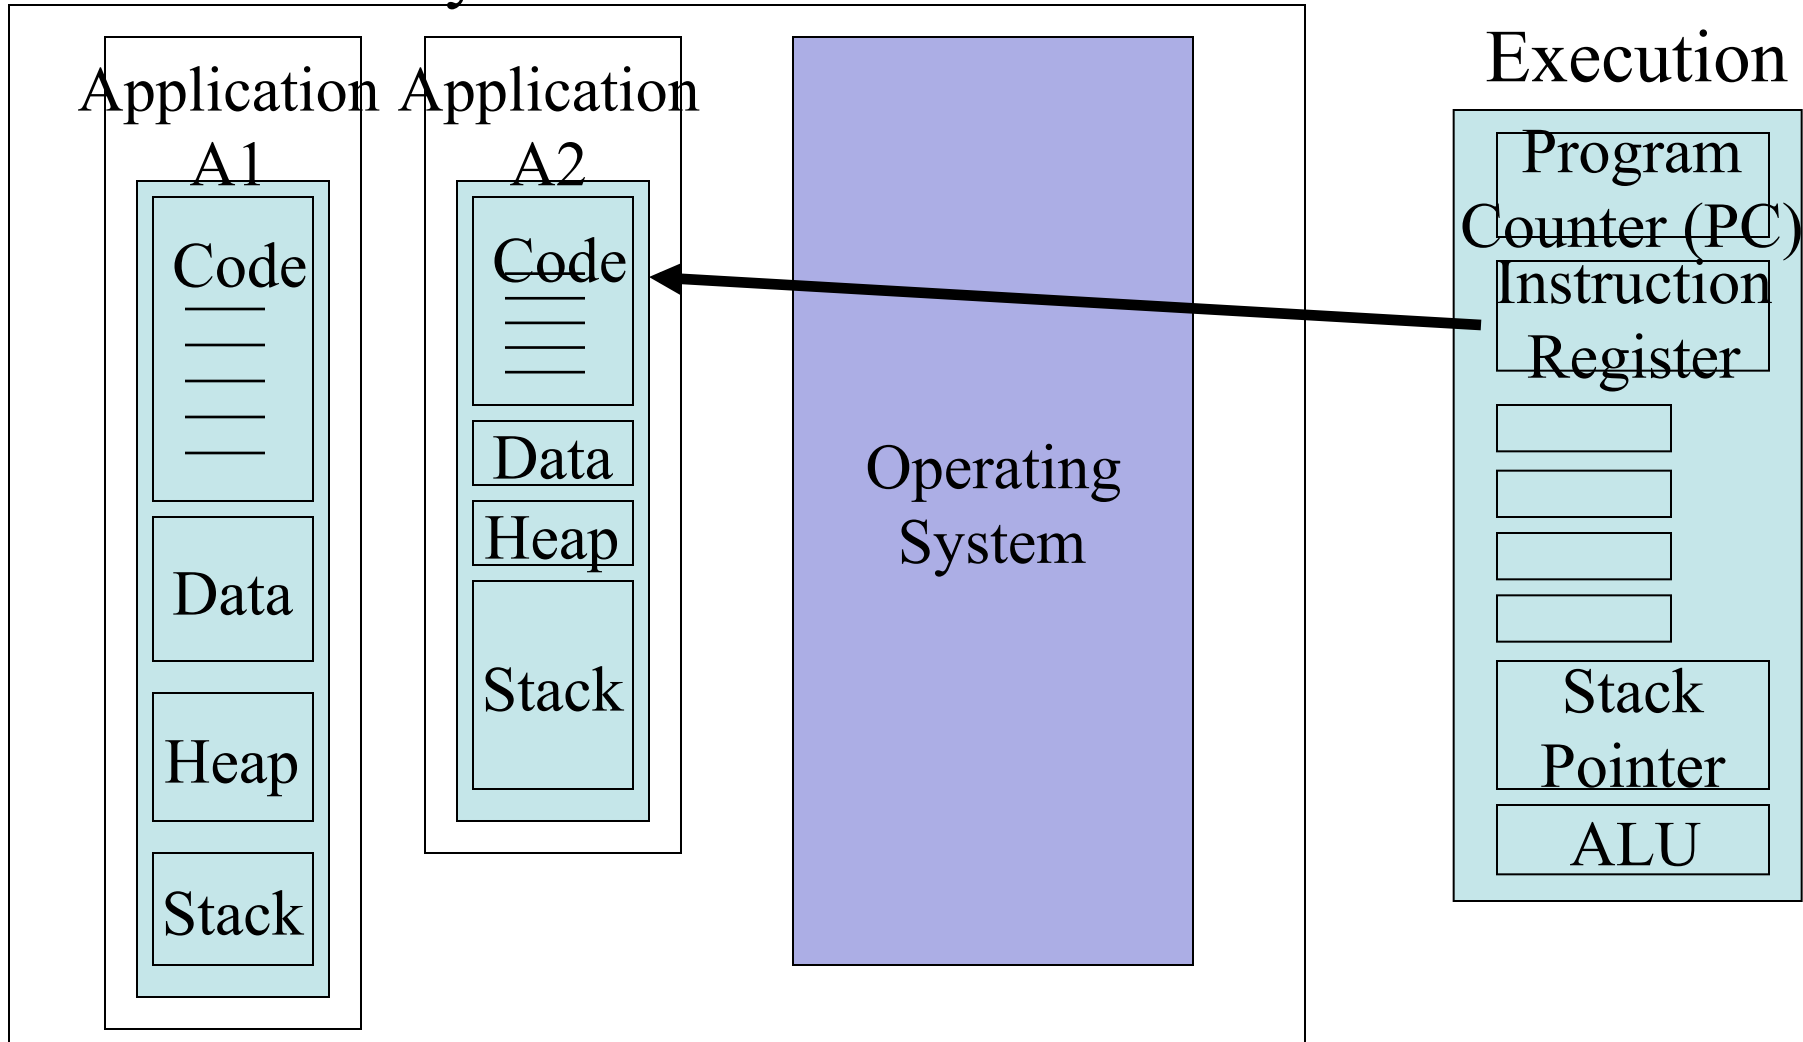


Multiple Applications + OS

Main Memory

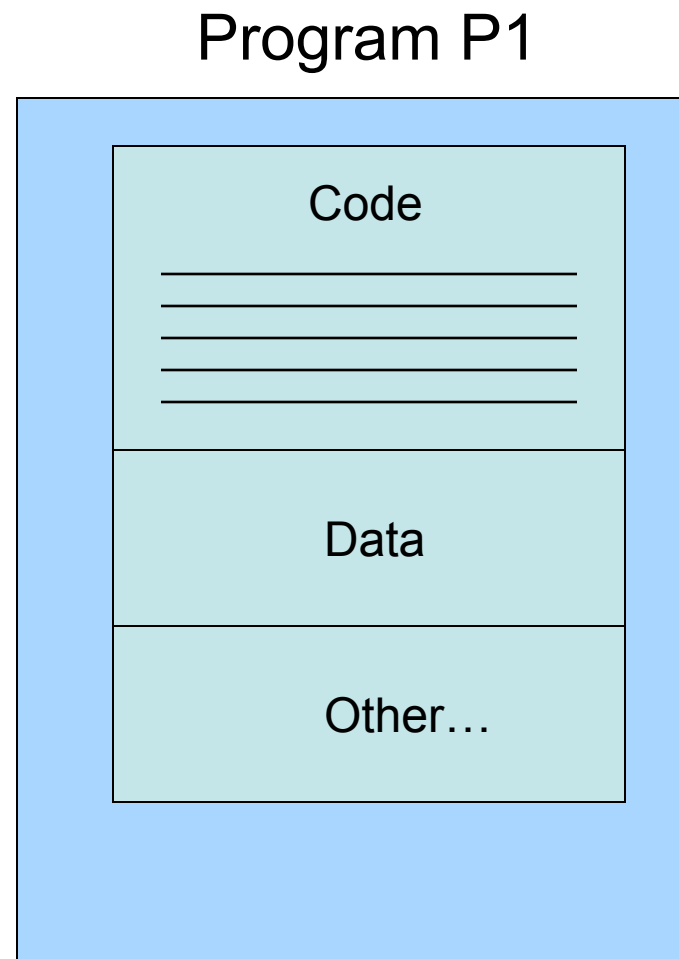
CPU

Execution

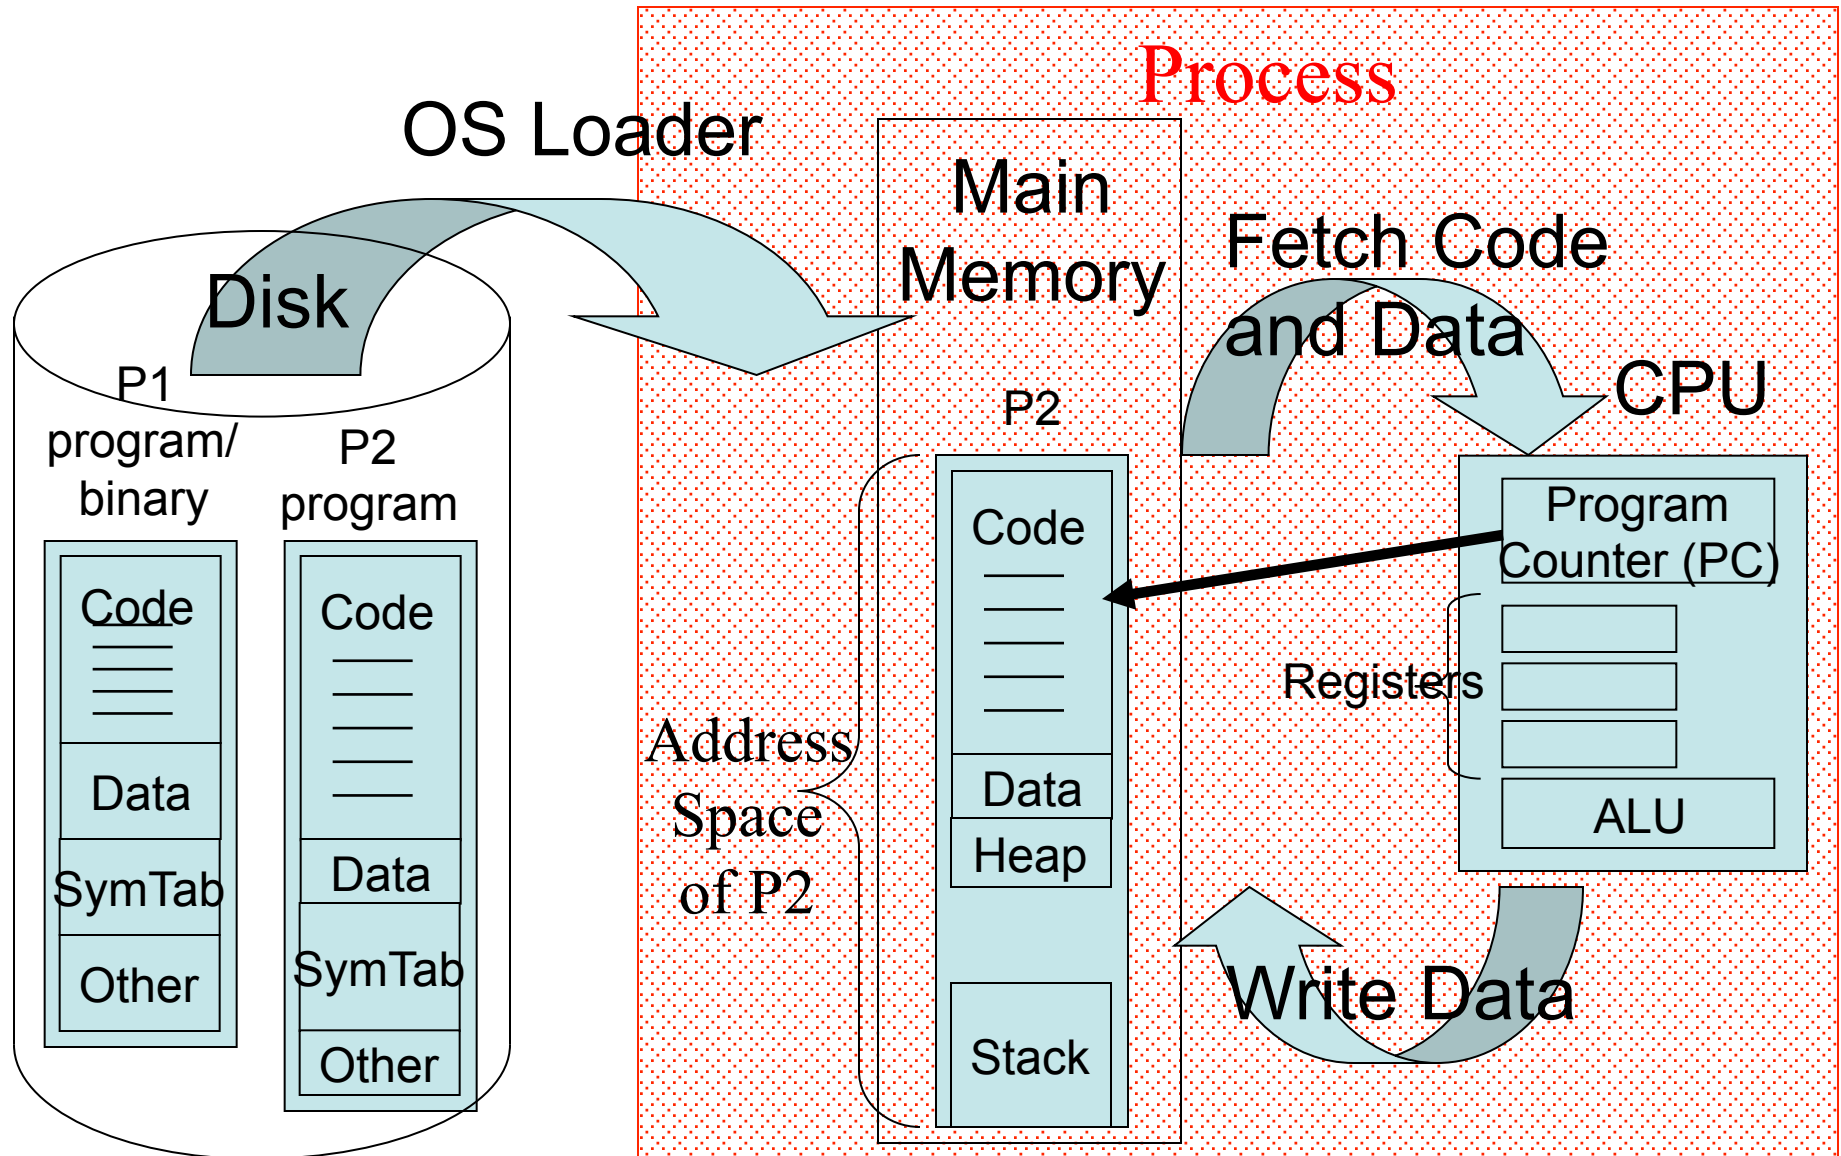


Chapter 3: What is a Process?

- A software *program* consist of a sequence of code instructions and data stored on disk
 - A program is a *passive* entity
- A *process* is a program *actively executing* from main memory within its *own address space*

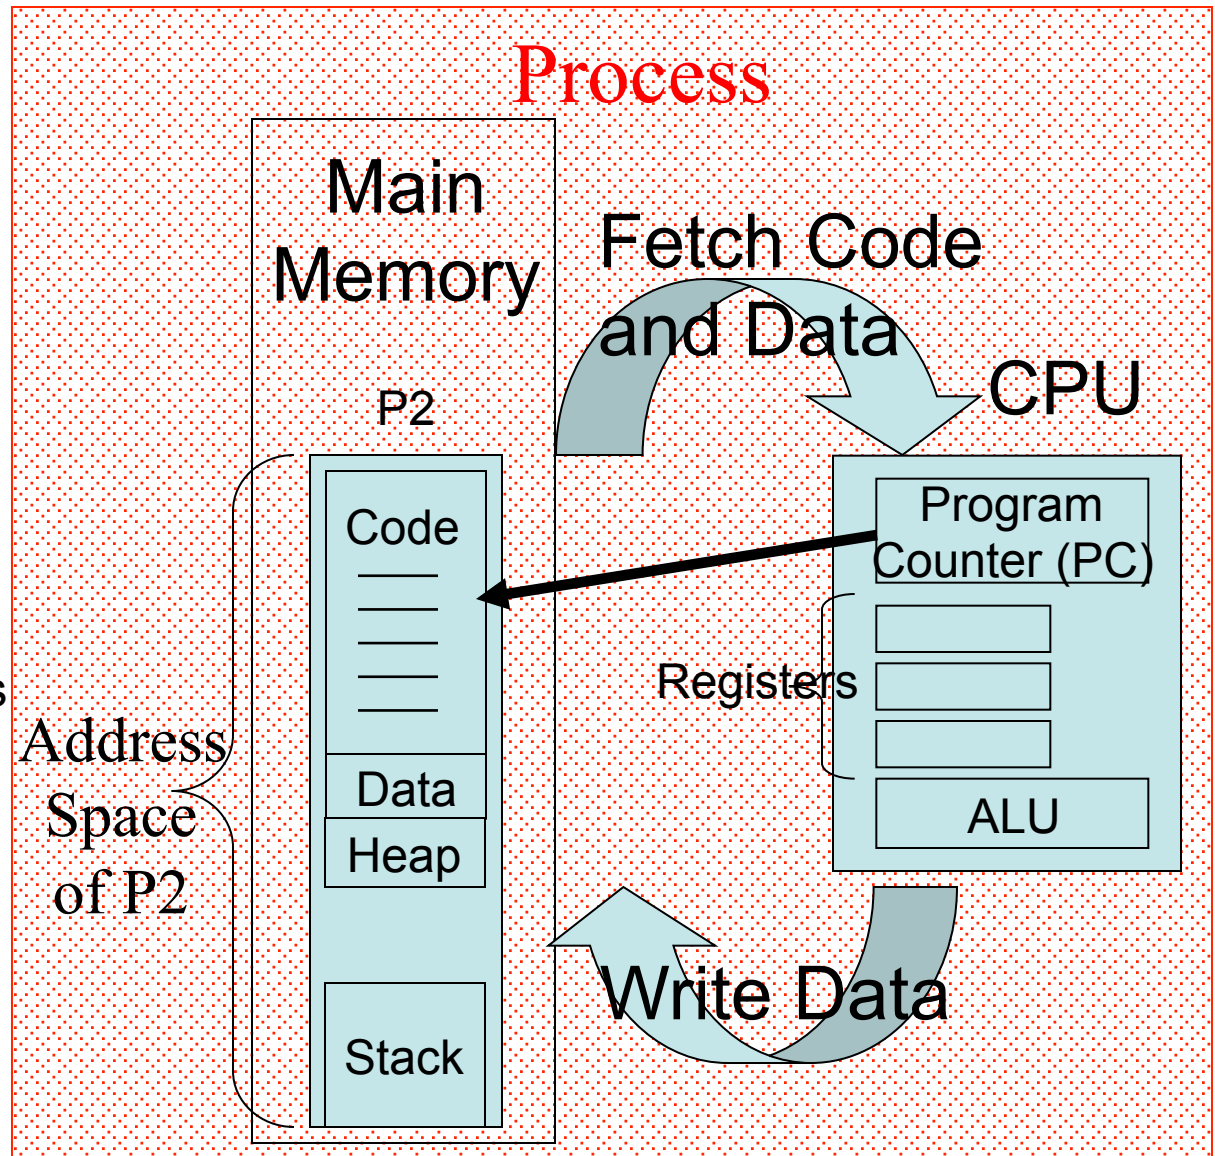


What Is a Process? (2)




What is a Process? (3)

- A *process* is a program *actively executing* from main memory
 - has a Program Counter (PC) and execution state associated with it
 - CPU registers keep state
 - OS keeps process state in memory
 - it's alive!
 - Owns its own *address space*
 - a limited set of (virtual) addresses that can be accessed by the executing code




What is a Process? (4)

- 2 processes may execute the same program code, but they are considered *separate execution sequences*
 - e.g. two shell terminals



A Process Executes in its Own Address Space

Main Memory

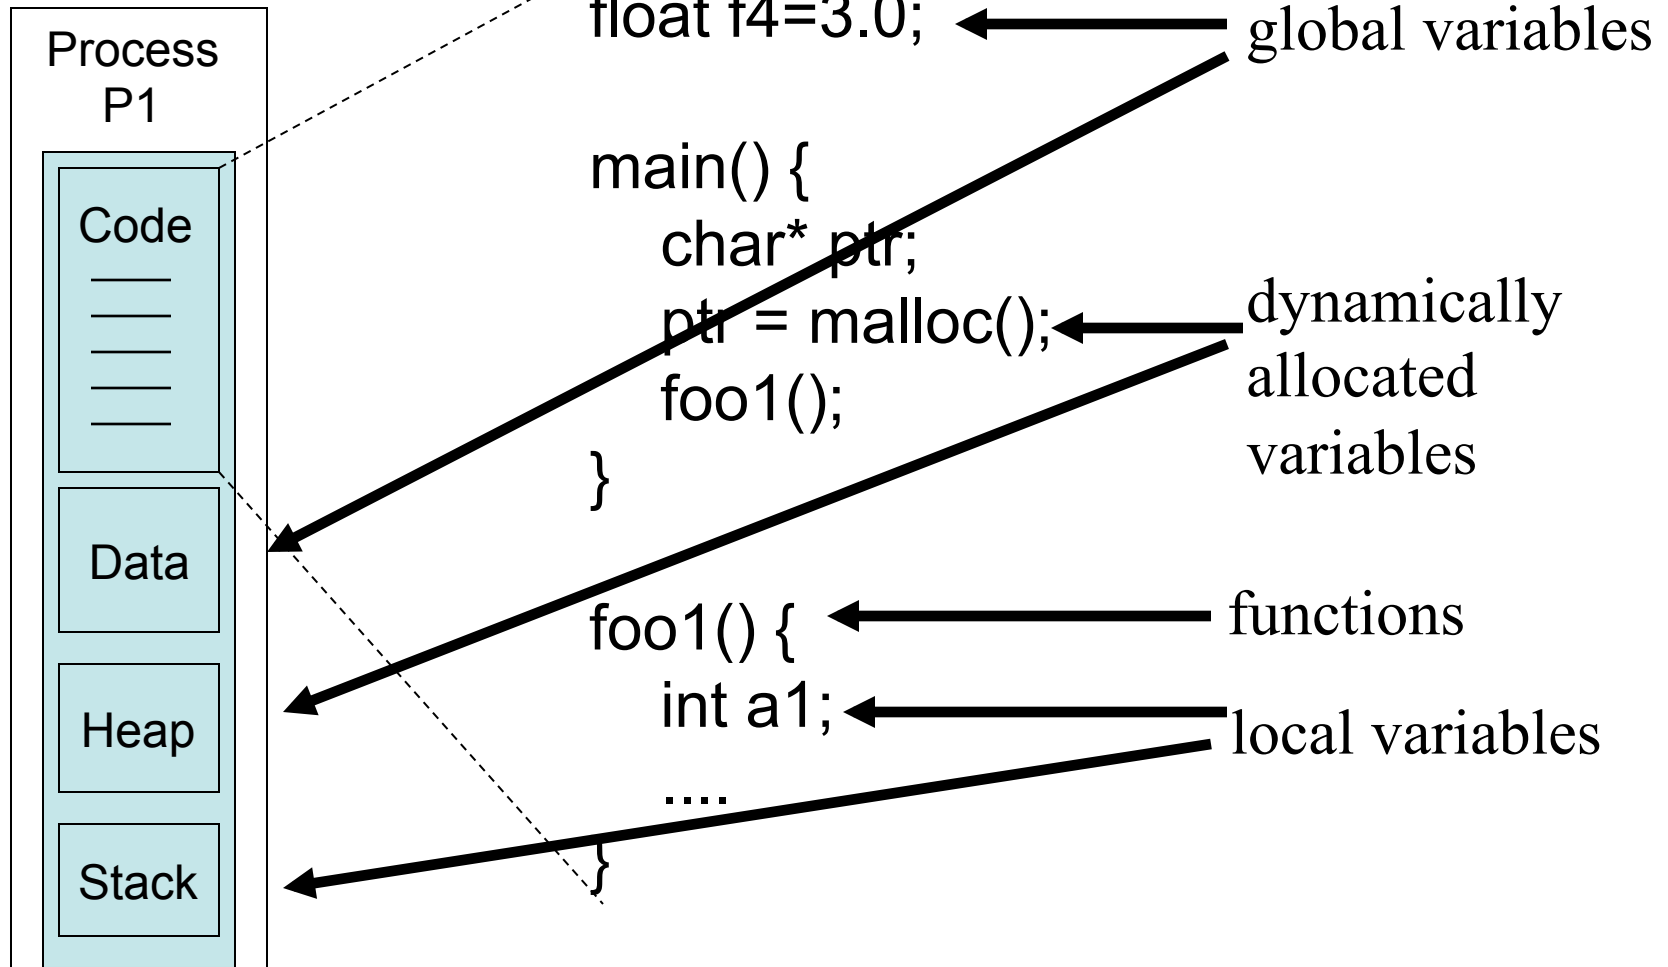


- OS tries to provide the illusion or *abstraction* to the process that it executes on its own abstract machine
 - in its own subset of RAM, i.e. its own address space – achieved using virtual memory paging
 - on its own subset (time slice) of the CPU – achieved by preemptive multitasking



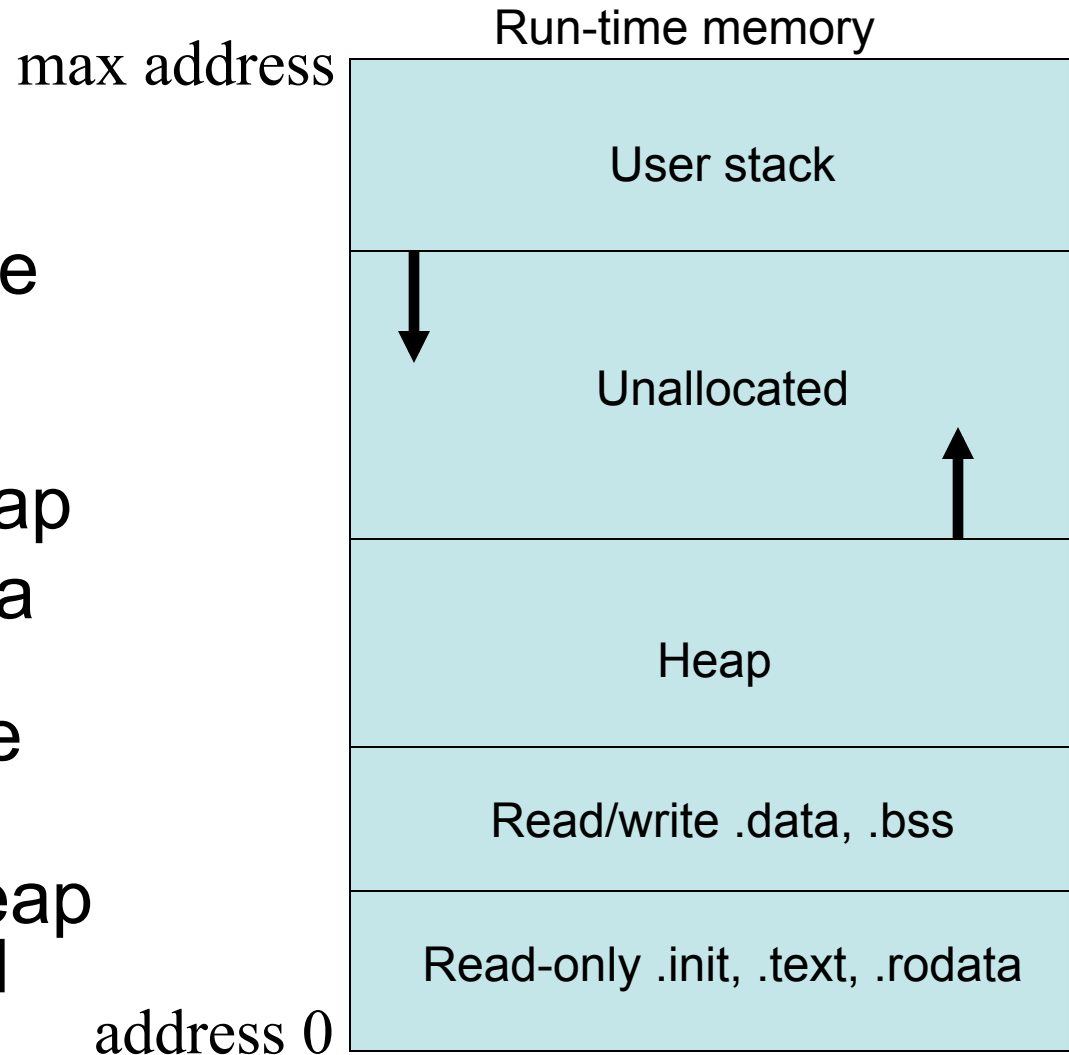
How is a Process Structured in Memory?

Main Memory



How is a Process Structured in Memory?

- Run-time memory image
- Essentially code, data, stack, and heap
- Code and data loaded from executable file
- Stack grows downward, heap grows upward



Applications and Processes

- Application = \sum Processes;
 - e.g. a server could be split into multiple processes, each one dedicated to a specific task (UI, computation, communication, etc.)
 - The Application's various processes talk to each other using Inter-Process Communication (IPC). We'll see various forms of IPC later.

